

CLEAR-MoE++ · Calibration-Driven Post-Training Expert Extraction from Pretrained Vision Transformers

DeiT-Small on Imagenette · $E = 4$ experts · 6 of 12 FFN layers expertised · cuBLAS dispatch: 243k tok/s

1 Input Images



Tench



Church



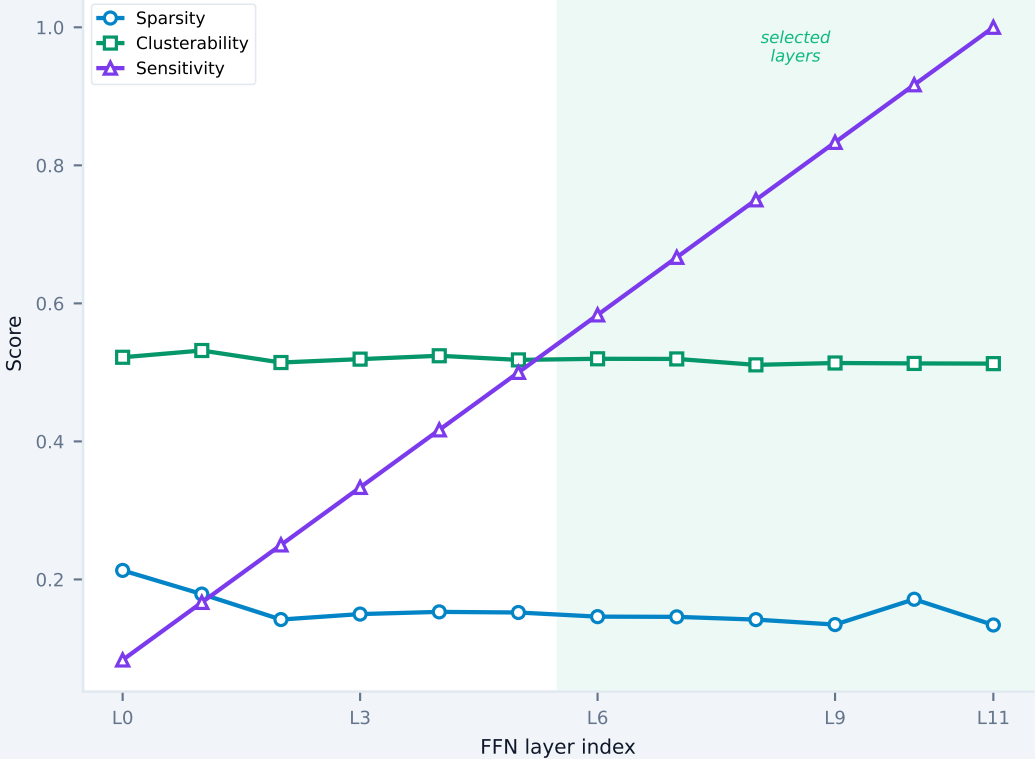
Golf Ball



Parachute

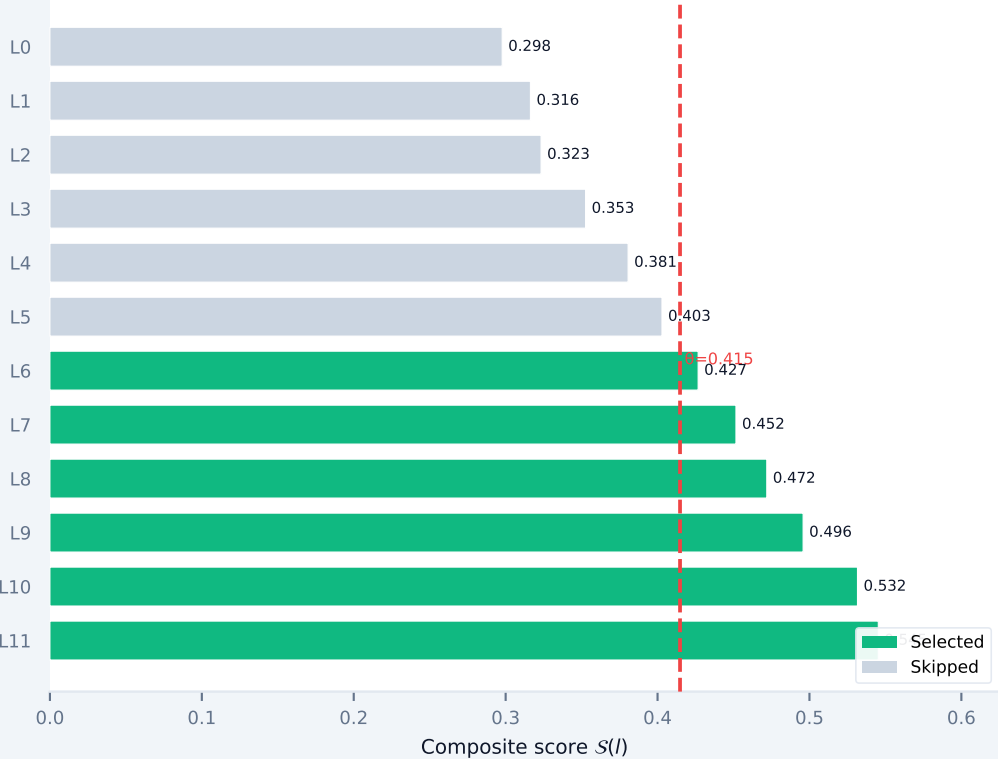
2 Calibration Pass

Activation statistics
(200 calib. images, DeiT-Small)



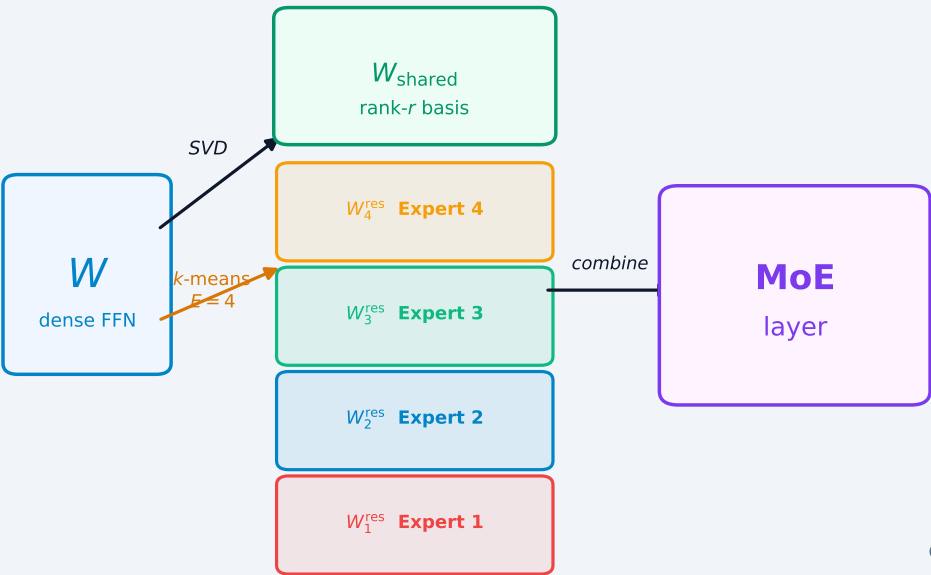
3 Layer Scoring

Layer scoring & selection
(last-half policy, L6-L11 selected)



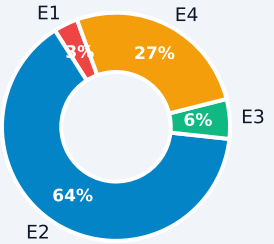
4 Expert Extraction

SVD + k -means decomposition
(recon. error 0.47–0.52 across L6–L11)

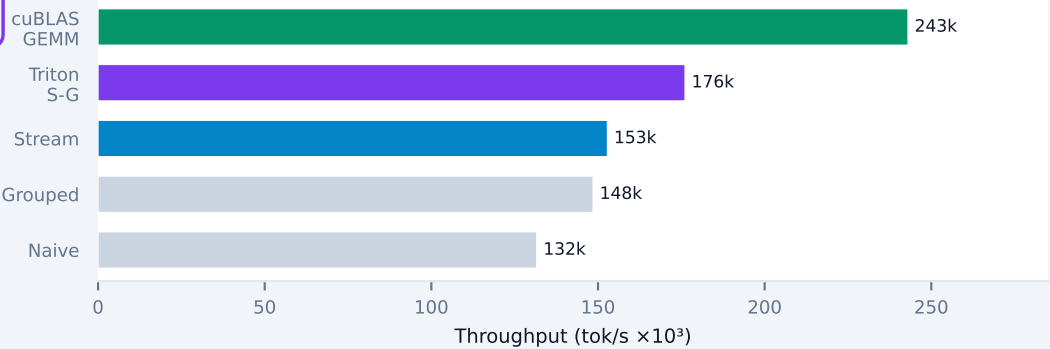


5 Router & Dispatch

Expert load distribution
(block L8, real routing stats)



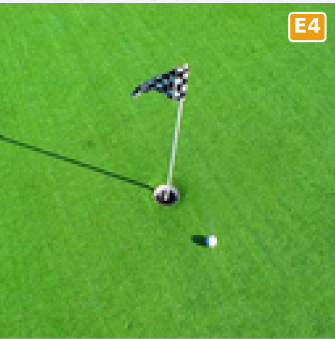
Dispatch backend throughput
(0% imbalance, $E = 4$, $d = 384$)



6 MoE Output



Tench



Golf Ball



Church



Parachute